

Watch Every ADCC Trials Finals From 2023-24 (Over 8 Hours Of Action!)

Video	https://www.youtube.com/watch?v=pwGbW5GZgfc
Title	Watch Every ADCC Trials Finals From 2023-24 (Over 8 Hours Of Action!)
Channel	FloGrappling
Uploaded	20240707
Video length	8:29:33
Frames sampled	12, one every 20s
Covers	8:25:29 to 8:29:09 of the video
Footage type	full event
Source note	ADCC Trials Finals 2023-24 compilation, unindexed remainder 8:25:29-8:29:33, coarse locator window

How to read this

Every frame carries a VIDEO-ABSOLUTE time -- seconds from the start of the video at the URL above, NOT seconds from the start of the bout -- printed as text directly above it. Those times run from 8:25:29 to 8:29:09.

Return every timestamp in that same video-absolute clock. If the bout starts partway into the video, say where it starts as a separate field -- do not rebase the event times to it. A bout-relative time reported as an absolute one places the whole fight in a different part of the broadcast, and the result looks plausible while being wrong.

THE CAPTION TEXT BESIDE EACH FRAME IS A GUIDE, NOT GROUND TRUTH. It is an auto-caption of live commentary: the commentator misnames athletes, talks about one bout while another is on screen, anticipates and speculates, and the text lags the picture. When the caption and the frame disagree, THE FRAME WINS. Use the caption to orient where to look and to break a tie on a label -- never to record an event you did not see in the frame itself.

For each event you can see, give:

<code>ts</code>	integer seconds, video-absolute
<code>label</code>	a technique name, taken verbatim from the label list
<code>actor</code>	which competitor did it, by name
<code>successful</code>	true if it landed, false if attempted and stopped
<code>type</code>	one of: control, submission, guard, takedown, pass, transition, sweep, escape
<code>points</code>	points the scoreboard awarded for THIS action; omit if you cannot read it

Points are the one thing the scoreboard is authoritative for, and reading footage is the first time they are visible at all -- a transcript only ever carried what a commentator said aloud. So: report `points` only when you can actually read the number change on the scoreboard between frames, attribute it to the action that caused it, and OMIT the field otherwise. Never write 0 to mean 'I could not tell' -- an action that scored nothing and an action nobody could score are different facts, and only an absent field holds the second. If the score jumps and you cannot see which action earned it, say so in `note` on the nearest event rather than guessing an attribution.

Also return, once per bout: the two competitors' names, the event/card name, the year, who won and how, `bout_start_seconds` -- the video-absolute second the bout begins -- and `final_score` as "a-b" if the scoreboard shows one. If this video contains more than one bout, return one such object per bout.

LABELS ARE A CLOSED VOCABULARY. Take each `label` verbatim from `node_library.json`, which lists every technique the library knows, most-used first. A label is normalised into a node key, and that key IS the node in every graph it reaches -- so a spelling the library does not have does not fail, it mints a second node, and one technique arrives split in two with no error anywhere. Where two listed names could both fit, take the one printed higher: that ordering is how often the corpus already uses it. If a frame genuinely shows something the list does not cover, use your own words and flag it as `new_label: true` rather than bending it into the nearest listed name.

NOT EVERY FRAME IS LIVE ACTION. A broadcast cuts to slow-motion replays, usually right after a finish or a scoring exchange, and a replay looks exactly like the thing it replays. Read one as new action and the same event is recorded twice, seconds apart, which downstream reads as a repeated technique rather than one shown twice. Tells: a graphics bug or logo appears, the clock stops or disappears, the camera angle changes, the action repeats something you already saw. Skip those frames rather than logging them, and say in the note which range you treated as replay.

AND NOT EVERY FRAME IS THIS BOUT. A broadcast cuts away -- to another mat, to a different match's scorecard, to a crowd shot or an ad -- and it may show two bouts in sequence with a gap between rounds. A reader found exactly that in this batch: a scorecard from an unrelated match sitting in the middle of the sampled frames. Events read off a cutaway land on the wrong athletes entirely, which is worse than missing them. Before logging any event, confirm the frame shows the two competitors you identified, and say in the note which ranges you excluded as not-this-bout.

THE SCOREBOARD'S NAME ORDER IS NOT A MAT POSITION. A lower-third graphic puts one name left and one right by broadcast convention; which competitor stands where on the mat is unrelated and changes shot to shot. Binding a name to a body through that order is the single highest-cost error available here, because `actor` is on every event -- get it backwards and the whole answer inverts at once, while still looking internally consistent. Bind identity to something carried ON the athlete instead: a name printed on the back of the rashguard, kit colour, sleeve length, an ankle band, hair. Then check that discriminator against a frame where the scoreboard and the bodies are visible together, and say in `identity\_verified\_by` which frame did it.

REPORT ONLY WHAT YOU ARE CONFIDENT OF -- but report uncertainty by getting COARSER, not by staying silent and not by guessing. If the frames show that something happened and you cannot tell exactly which technique it was, use the generic label the library already carries -- `Guard Pass` rather than a named pass, `Takedown` rather than a named entry, `Submission` rather than a named finish, and likewise `Sweep`, `Escape`, `Guard`. A generic label is the honest answer at the resolution these frames give you; a specific one you inferred is a claim about a technique nobody saw, and once it is in the athlete's graph nothing downstream can tell it from an observed one.

This matters most across the sampling gap. At one frame every 20s a technique can start and finish between two samples, leaving only its result visible -- a competitor who was on top now now underneath, a grip that is now a finish. Recording that result generically is correct. Naming the entry you did not see is not, however obvious it feels. The same rule covers `actor`: if you cannot tell which competitor did it, say so in `note` rather than picking one, because a right technique on the wrong athlete is two errors.

The result line at the top, when there is one, comes from the athlete's BJJ Heroes record. It is there so you can tell the two competitors apart and know the bout ended -- it is NOT evidence of anything you cannot see. Do not report a finish, a technique or a score because that line names it: an armbar listed as the method and an armbar visible in the frames are two different claims, and only the second belongs in your answer.

Only report what a frame actually shows. The scoreboard settles POINTS and nothing else: a technique you infer from a score change but never see is worth less than an omission, because a wrong label propagates into the athlete's graph and their rating and nothing downstream can tell it from an observed one. An omission is recoverable; an invented event is not.

8:25:29 (30329s)



Don't do nothing. Go.

8:25:49 (30349s)



(no narration in this window)

8:26:09 (30369s)



[Music]

8:26:29 (30389s)



(no narration in this window)

8:26:49 (30409s)



We out. We out.

8:27:09 (30429s)



[Applause] [Music] [Applause] [Music]

8:27:29 (30449s)



Back. Back. Let's [Applause] [Music]

8:27:49 (30469s)



go [Music] down.

8:28:09 (30489s)



[Music] She's all right. [Applause]

8:28:29 (30509s)



[Applause] It go Nice.

8:28:49 (30529s)



(no narration in this window)

8:29:09 (30549s)



Yep. [Applause]